



Unpaired Image-to-Image Translation Using CycleGAN

HINAL SHAH

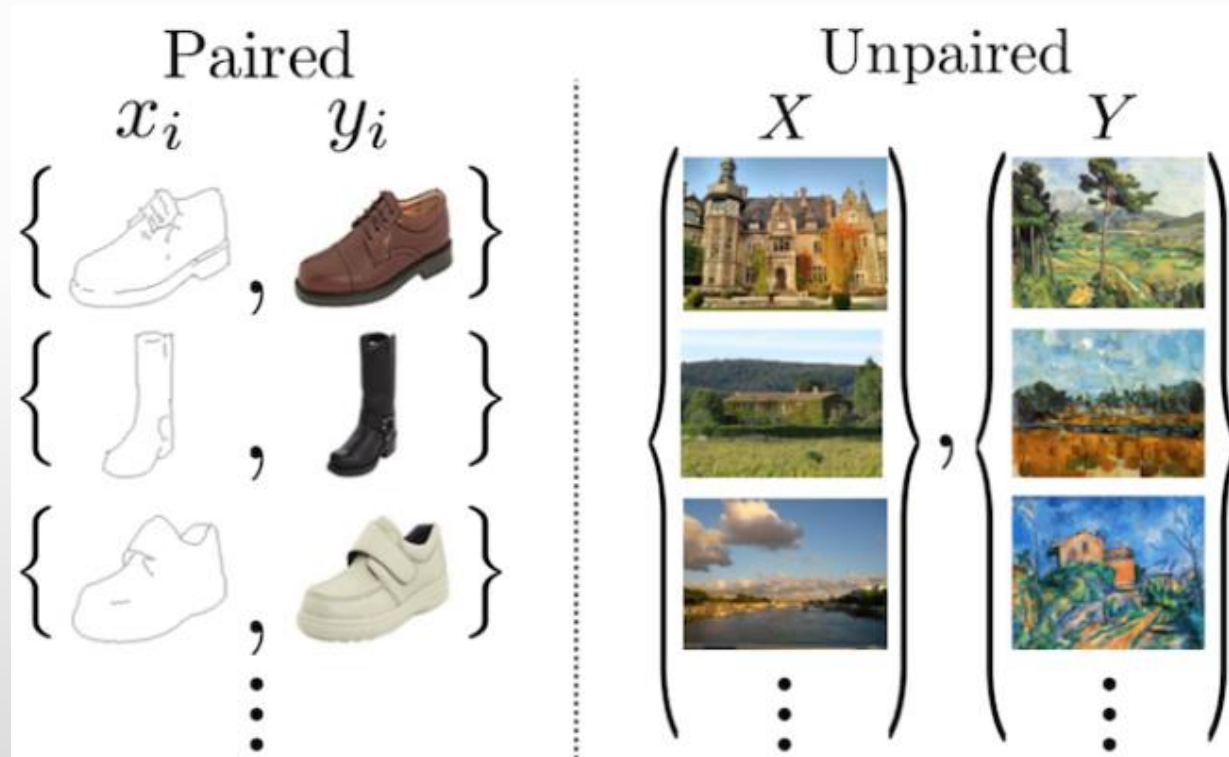
Introduction

- Image-to-image translation is the task of transforming an image from one domain to another by learning mapping between an input image and an output image



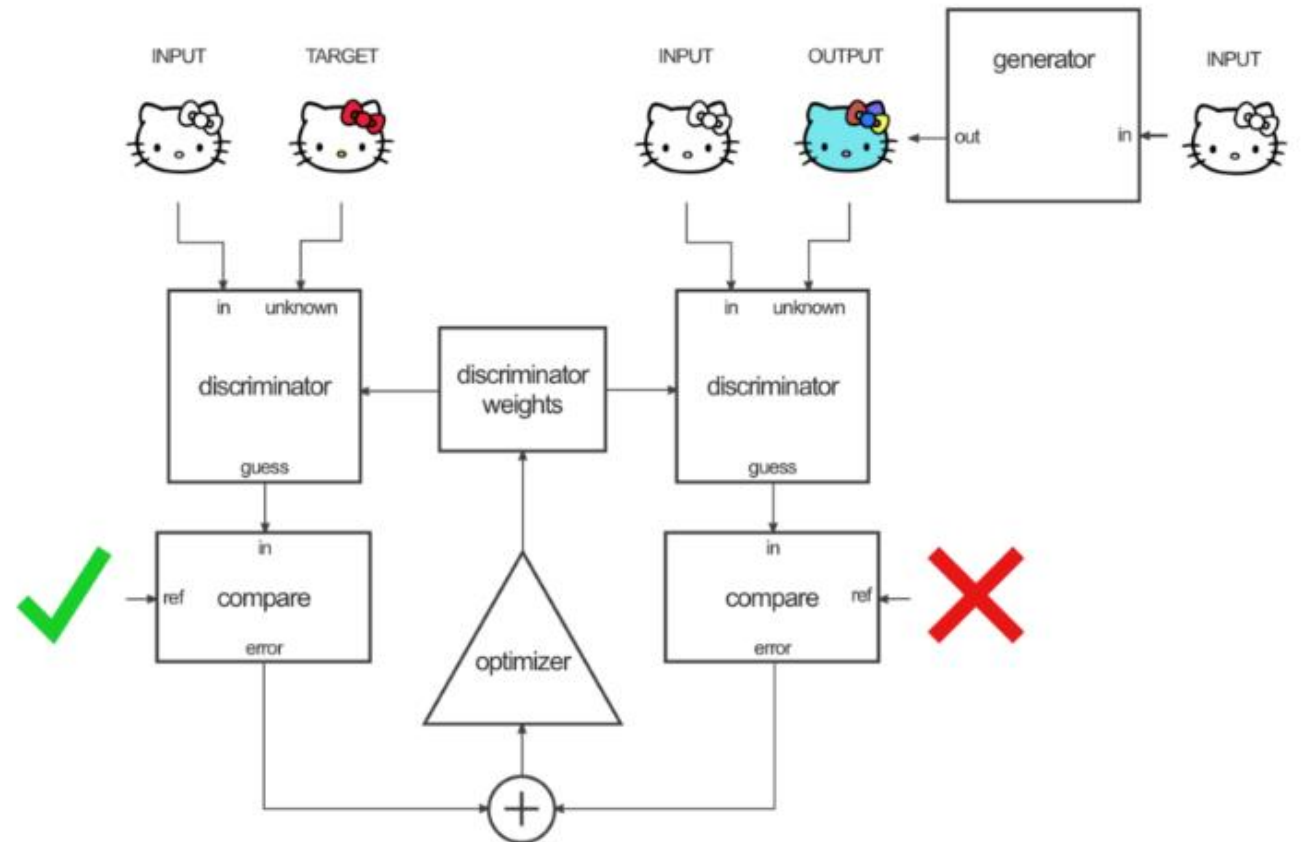
Introduction

- Paired data is harder to find in most domains, and not even possible in some, so the unsupervised training capabilities of CycleGAN is useful.



Pix2Pix Overview

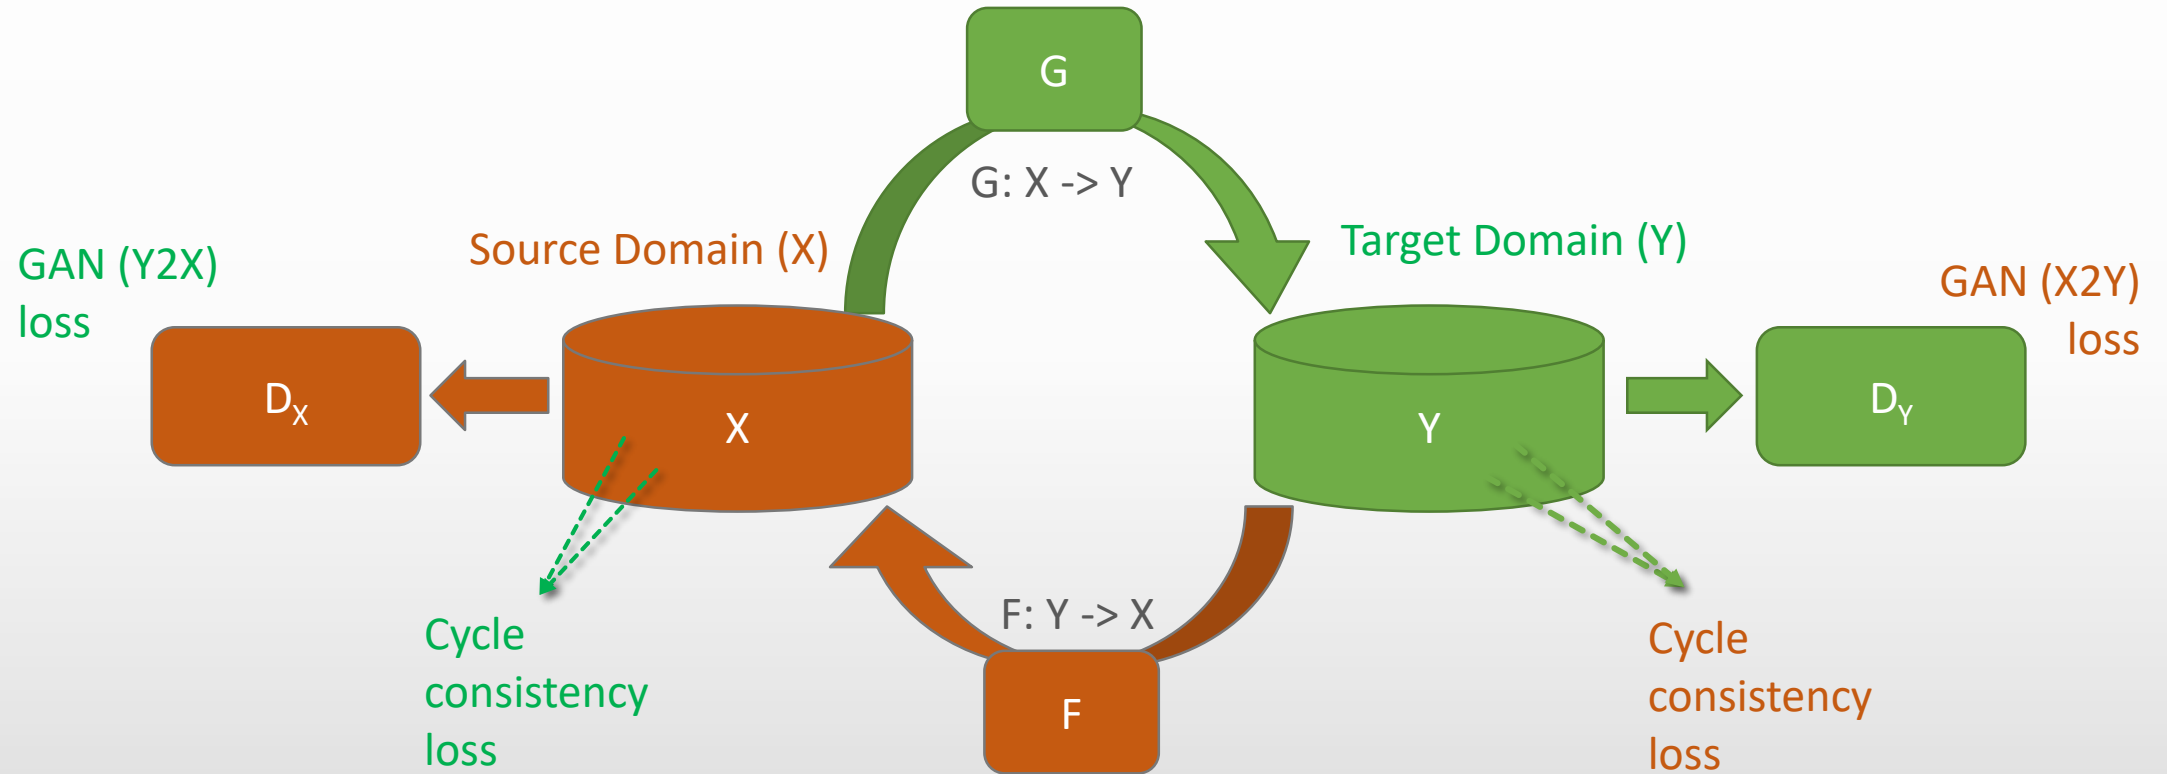
In paired dataset, input and output images share some common features. So there is a meaningful transformation
(Eg: Pix2Pix model)




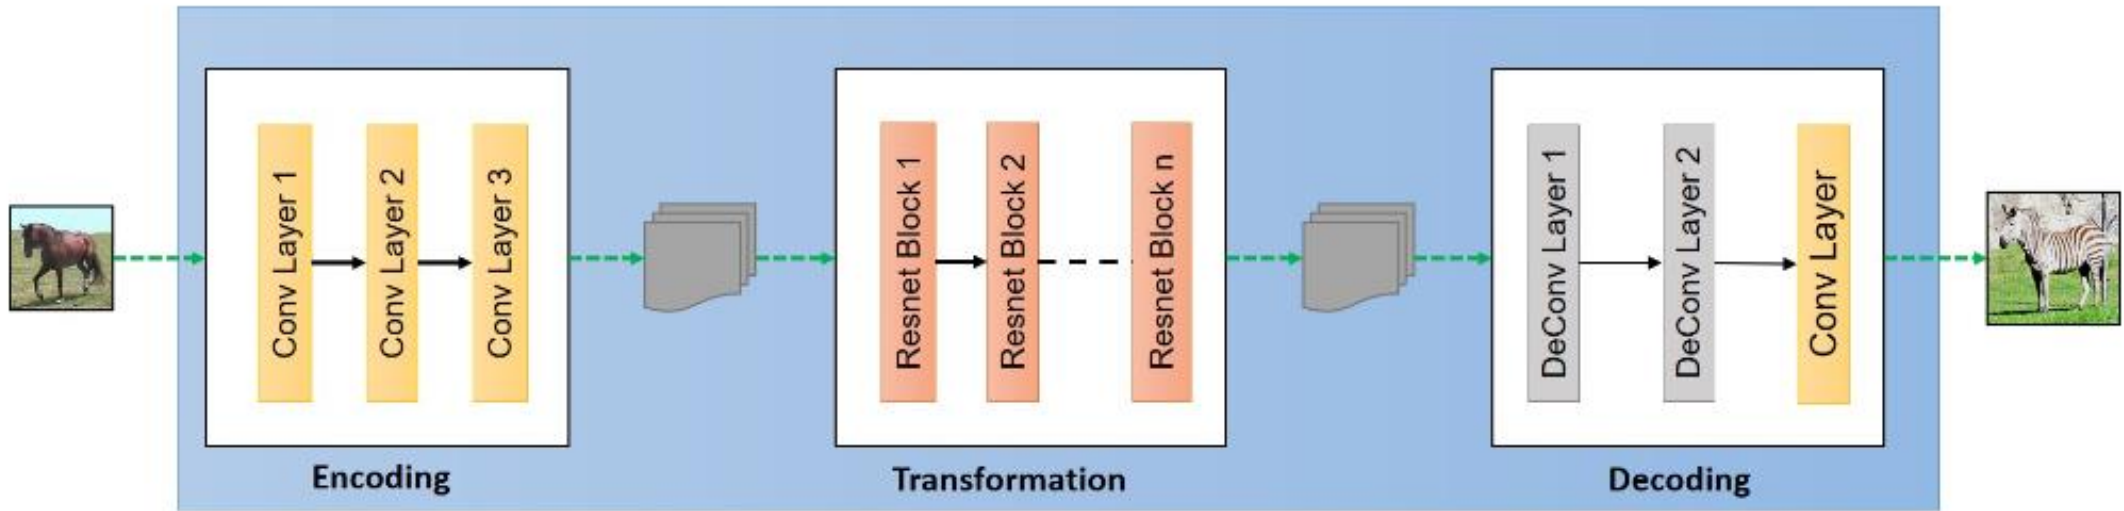
CycleGAN

- In cycleGAN, datasets are unpaired. So, to have meaningful mapping between source and domain images, they must share some feature that can be used to map this output image back to input image
- And this can be done by using another generator that must be able to map back this output image back to original input image.
- Since mapping is done in both direction Source \rightarrow Target Domain and Target \rightarrow Source Domain, we have two Discriminators along with two Generators.

CycleGAN Architecture


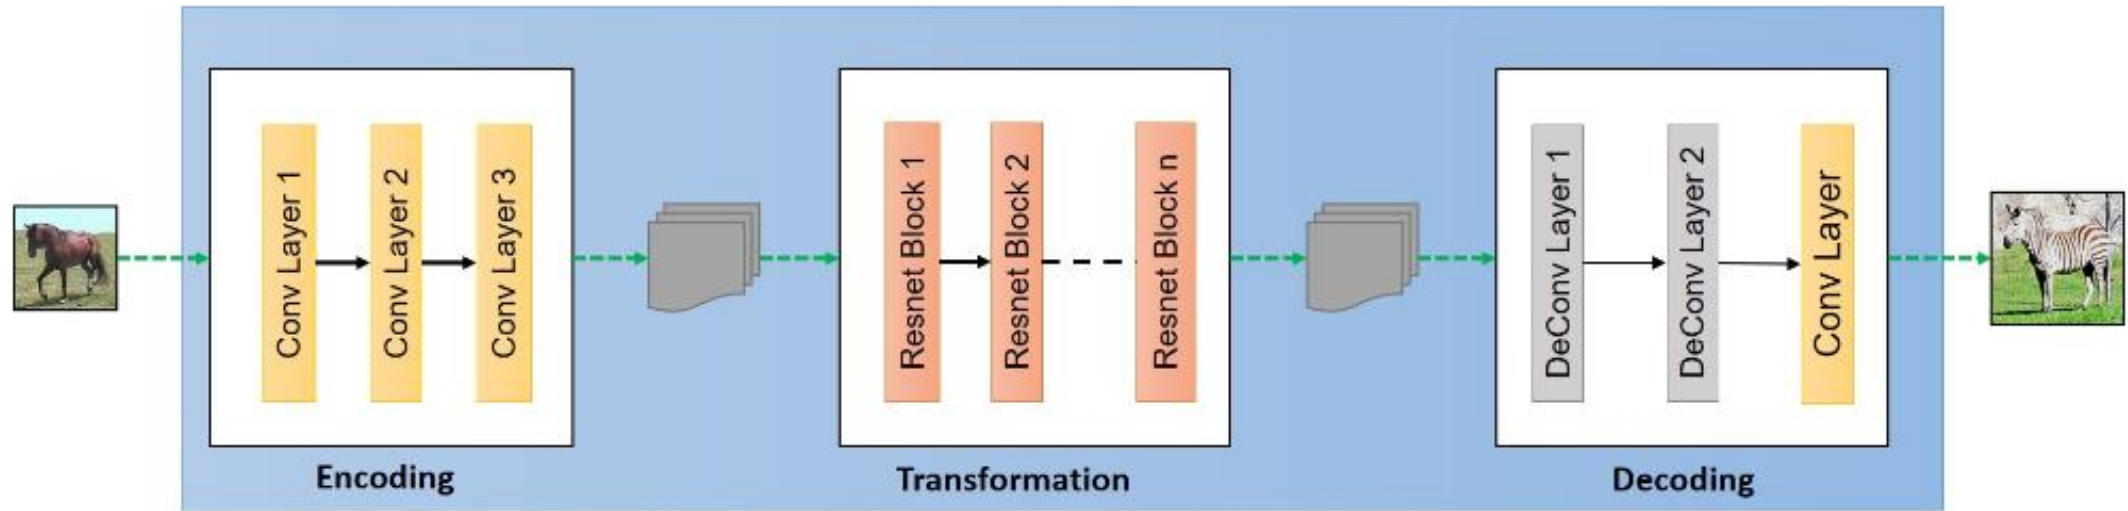


Generator Architecture



6 Resnet Blocks for 128 x 128 images
Conv Layers: 32, 64, 128,
Resnet Layers: R128, R128, R128, R128,
R128, R128,
DeConv Layers: 64, 32, 3

Generator Architecture




9 Resnet Blocks for 256 x 256 and higher resolution images

Conv Layers: 64, 128, 256

Resnet Layers: R256, R256, R256, R256, R256, R256, R256, R256, R256

DeConv Layers: 128, 64, 3

Discriminator Architecture



For **128 x 128** and higher resolution images

Conv Layers: 32, 64, 128, 256

Last Conv Layer: 1 (sigmoid layer)

For **256 x 256** and higher resolution images

Conv Layers: 64, 128, 256, 512

Last Conv Layer: 1 (sigmoid layer)

Objective Function

Adversarial Loss ($G: X \rightarrow Y$):

$$\mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))]$$

Cycle Consistency Loss:

$$\mathcal{L}_{\text{cyc}}(G, F) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1]$$

Full Objective Function:

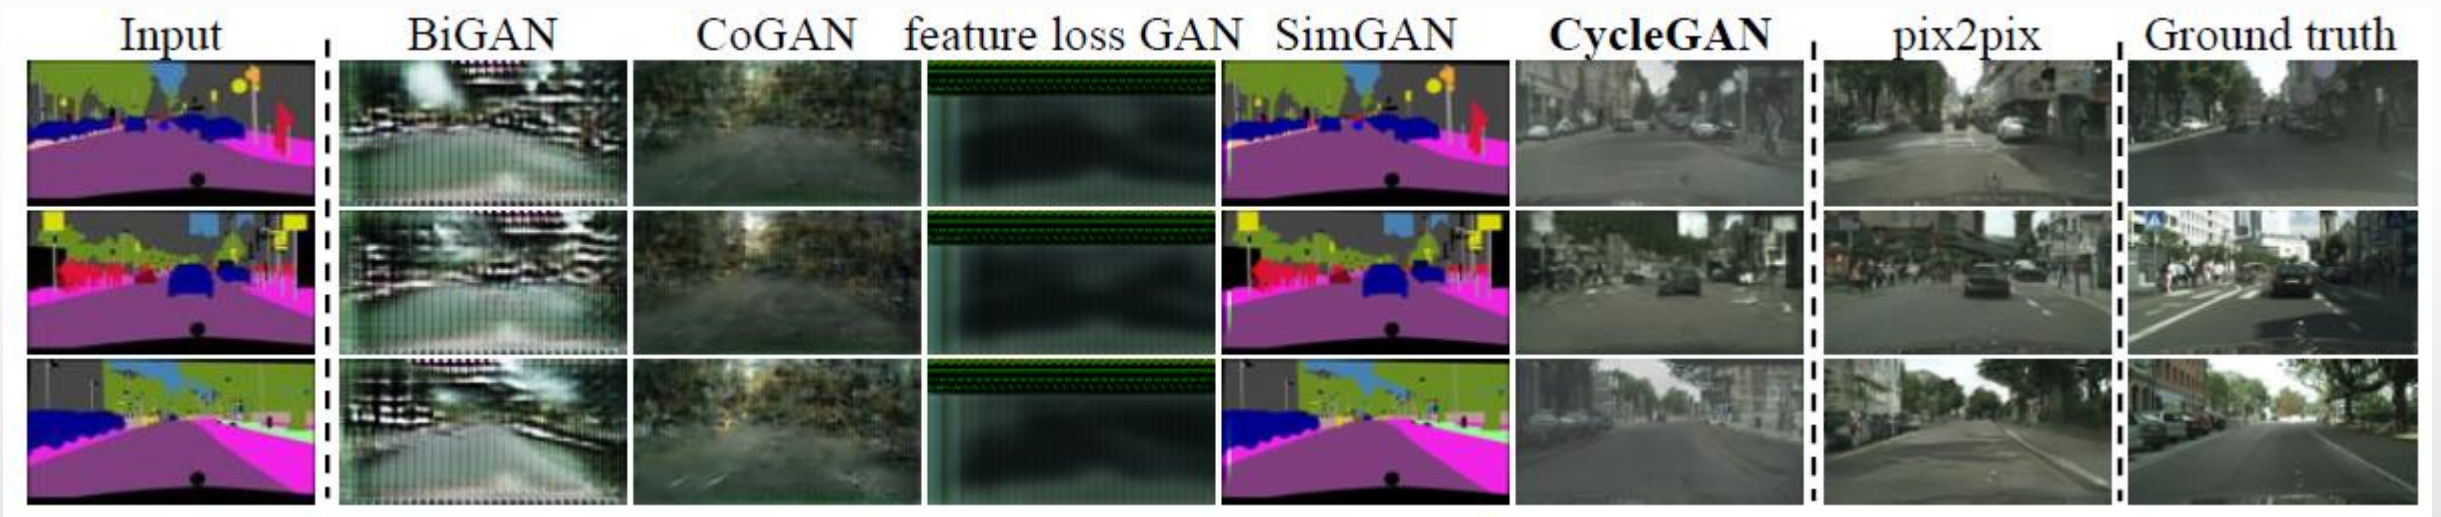
$$\begin{aligned} \mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) \\ & + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{\text{cyc}}(G, F), \end{aligned}$$

So we are trying to solve:

$$G^*, F^* = \arg \min_{G, F} \max_{D_X, D_Y} \mathcal{L}(G, F, D_X, D_Y)$$

Evaluation

- No ground Truth values available as datasets are unpaired



Evaluation on Cityscapes Dataset

Results

Input



Output





Sunflower → Daisy

Input



Output



Daisy → Sunflower

Results

Input



Output






Horse → Zebra

Input



Output



Zebra → Horse

Results

Input



Output



Face → Cartoon

Input



Output



Cartoon → Face

Results

Input




Output



Input




Output



Input




Output



Some Failed Cases

Conclusion

This method works well in performing tasks that involve color and texture changes like photo enhancement



SmartPhone Snaps → DSLR Photographs

Conclusion

This method works well in performing tasks that involve color and texture changes like photo enhancement



Orange → Apple



Summer Yosemite → Winter Yosemite

Conclusion

It does not work well in tasks that require geometric changes




Cat → Dog



Dog → Cat

Conclusion

- Does not perform well in translating images that has complicated backgrounds
- It does only one-to-one mapping



Horse → Zebra

Questions?
